

Activations Functions :

What is an Activation Function?

An activation function in a neural network is a mathematical function applied to each neuron's output (also called the neuron's activation). It determines whether a neuron should be activated or not, introducing non-linearity into the model, which allows the network to learn complex patterns.

Why Do We Need Activation Functions?

1. Introduce Non-linearity:

- Without non-linearity, a neural network, no matter how many layers it has, would behave like a single-layer linear model.
- Non-linear activation functions allow the network to learn and represent complex patterns and functions.

2. Control the Output:

- Activation functions can control the range of the output values (e.g., sigmoid outputs between 0 and 1).
- This is particularly useful for tasks like classification where outputs need to represent probabilities.

3. Gradient Flow:

- Proper activation functions ensure gradients flow properly during backpropagation.
- This helps in updating weights effectively and avoiding issues like vanishing gradients.

Good Features in Activation Functions

1. Non-linearity:

- Enables the network to solve non-trivial problems.
- Example: ReLU introduces non-linearity by zeroing out negative values.

2. Differentiability:

- Essential for backpropagation to compute gradients.

- Example: Sigmoid and Tanh functions are differentiable at all points.

3. **Computational Efficiency:**

- Should not add significant computational overhead.
- Example: ReLU is computationally efficient as it only involves a simple threshold operation.

4. **Avoiding Vanishing/Exploding Gradients:**

- Ensures stable training by maintaining gradient magnitudes.
- Example: ReLU helps mitigate the vanishing gradient problem.

5. **Zero-Centered Outputs:**

- Helps in faster convergence by making the output range symmetric around zero.
- Example: Tanh outputs range from -1 to 1.

6. **Sparse Activation:**

- Only a subset of neurons should be active at a time, leading to efficient computation.
- Example: ReLU only activates neurons with positive inputs.

Bad Features in Activation Functions

1. **Non-differentiability:**

- Prevents the use of gradient-based optimization.
- Example: Binary Step Function is not differentiable.

2. **Vanishing Gradients:**

- Leads to very small gradients, making it hard for the network to learn.
- Example: Sigmoid can cause vanishing gradients for large positive or negative inputs.

3. **Exploding Gradients:**

- Causes gradients to become excessively large, leading to unstable training.
- Example: Functions that lead to large outputs for large inputs can cause exploding gradients.

4. Dead Neurons:

- Neurons that never activate and thus do not contribute to learning.
- Example: ReLU can cause dead neurons when inputs are negative.


5. Not Zero-Centered:

- Slows down convergence as gradients may consistently move in a certain direction.
- Example: Sigmoid outputs are not zero-centered (range from 0 to 1).

Activation Functions in Neural Networks: Detailed Analysis

1. Binary Step Function

- **Formula:**




- **Range:** {0,1}
- **Usage:** Rarely used in practice due to its simplicity and limitations. Sometimes used in perceptrons for binary classification.
- **Advantages:** Simple implementation.
- **Disadvantages:**

- Not useful for multi-class classification.
- Gradient is zero everywhere, hindering backpropagation.
- Not differentiable, making it unsuitable for gradient-based optimization.

2. Linear Function

- **Formula:**


-



- **Range:** $(-\infty, \infty)$
- **Usage:** Generally used in simple regression tasks or in the final layer of a neural network for regression problems.
- **Advantages:**
 - Simple and interpretable.
 - No vanishing gradient problem.
- **Disadvantages:**
 - Gradient is constant and does not depend on x .
 - Poor at capturing complex patterns.
 - The network may not learn non-linear boundaries.

3. Sigmoid Activation Function


- **Formula:**



- **Range:** (0,1)
(0,1)(0, 1)
- **Usage:** Commonly used in binary classification problems and in the output layer of binary classifiers.
- **Advantages:**
 - Smooth and differentiable.
 - Output can be interpreted as probabilities.
- **Disadvantages:**
 - Vanishing gradient problem for large positive or negative values.
 - Output is not zero-centered, which can slow down convergence.

4. Tanh (Hyperbolic Tangent) Function


- **Formula:**



- **Range:** $(-1,1)$
 $(-1,1)$
- **Usage:** Often preferred over sigmoid in hidden layers due to its zero-centered output.
- **Advantages:**
 - Zero-centered output.
 - Smooth and differentiable.
- **Disadvantages:**
 - Vanishing gradient problem similar to sigmoid.
 - Computationally more expensive than ReLU.

5. ReLU (Rectified Linear Unit)

- **Formula:**




- **Range:** $[0, \infty)$
 $[0, \infty)[0, \infty)$
- **Usage:** Widely used in hidden layers of neural networks.
- **Advantages:**
 - Computationally efficient.
 - Helps mitigate the vanishing gradient problem.
 - Sparse activation (neurons are activated only when necessary).
- **Disadvantages:**
 - Can create "dead neurons" (neurons that never activate if $x < 0$ consistently).

6. Leaky ReLU

- **Formula:**
-


Leaky ReLU Activation Function



- **Range:** $(-\infty, \infty)$
 $(-\infty, \infty)$
- **Usage:** Used to address the dead neuron problem in ReLU.
- **Advantages:**
 - Mitigates the dead neuron problem by allowing a small gradient when $x < 0$.
- **Disadvantages:**
 - The choice of the leakage factor (0.01) may not be optimal for all tasks.

7. Parametric ReLU

- **Formula:**



- **Range:** $(-\infty, \infty)$
 $(-\infty, \infty)$
 - **Usage:** Used to address dead neurons with a learnable parameter a .
- aa

- **Advantages:**

- Allows the network to learn the optimal value of α for better performance.


α

- **Disadvantages:**

- More complex than ReLU and Leaky ReLU.

8. Exponential Linear Unit (ELU)

- **Formula:**



- **Range:** $(-\infty, \infty)$ for $x < 0$ and $[0, \infty)$ for $x \geq 0$

$(-\infty, \infty)$

$x < 0$

- **Usage:** Used to avoid dead neurons and to have smooth gradients for negative inputs.

- **Advantages:**

- Helps mitigate the vanishing gradient problem.
- The negative saturation (when $x < 0$) helps regularize the network.

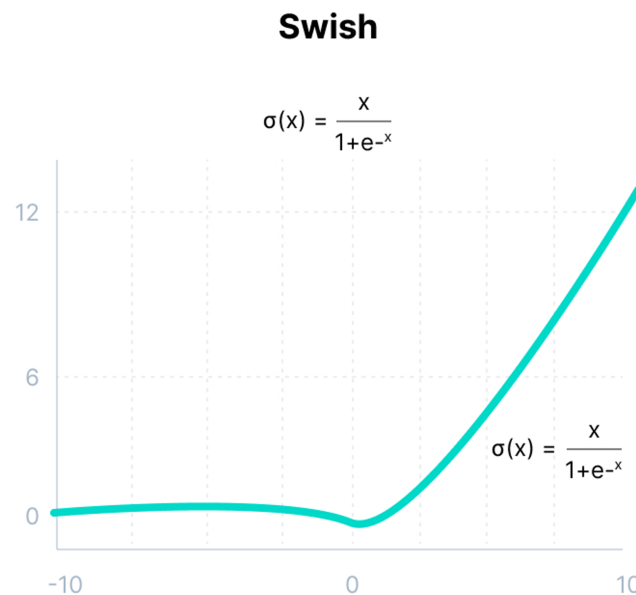
$x < 0$

- **Disadvantages:**

- More computationally intensive due to the exponential function.

9. Swish


- **Formula:**



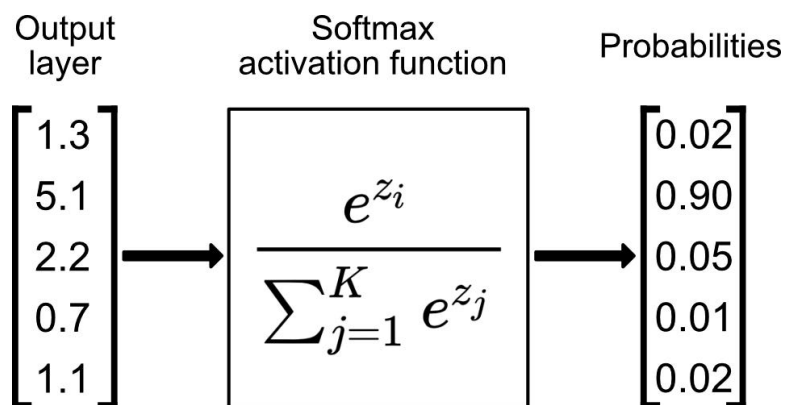
- **Range:** $(-\infty, \infty)$
 $(-\infty, \infty)$
- **Usage:** Used in deeper neural networks and has shown to outperform ReLU in certain tasks.
- **Advantages:**
 - Smooth and differentiable.
 - No dead neuron problem.
 - Better performance in deep networks.
- **Disadvantages:**
 - More computationally intensive than ReLU.
 - Non-monotonic, which can complicate understanding of the activation's effect.

10. Softmax

- **Formula:**



•



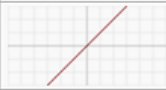








•

- **Range:** [0,1] for each output neuron, and the sum of all outputs is 1.
[0,1][0, 1]
- **Usage:** Used in the output layer for multi-class classification problems.
- **Advantages:**
 - Converts logits to probabilities, making them interpretable.
 - Ensures that outputs sum to 1, representing a valid probability distribution.
- **Disadvantages:**

- Computationally intensive due to exponentiation and normalization.
- Can suffer from the vanishing gradient problem if logits are large.

Summary of Usage and Recommendations

- **Binary Step Function:** Simple tasks with binary outputs, though rarely used in practice due to its limitations.
 - **Linear Function:** Regression problems, particularly in the output layer.
 - **Sigmoid Function:** Binary classification, particularly in the output layer.
 - **Tanh Function:** Preferred over sigmoid for hidden layers due to zero-centered output.
 - **ReLU Function:** Default choice for hidden layers in most networks.
 - **Leaky ReLU:** Used when encountering dead neurons with ReLU.
 - **Parametric ReLU:** Advanced variant of Leaky ReLU when even finer control over negative slopes is needed.
 - **ELU:** Used when the benefits of ReLU are desired, but with smoother negative saturation to avoid dead neurons.
 - **Swish:** Used in deeper networks for better performance, though computationally more demanding.
 - **Softmax:** Used in the output layer for multi-class classification to convert logits to probabilities.
-
- **Binary Classification:** Sigmoid (output layer), Tanh/ReLU (hidden layers).
 - **Multi-class Classification:** Softmax (output layer), ReLU (hidden layers).
 - **Regression:** Linear (output layer), ReLU (hidden layers).
 - **General Use:** ReLU for hidden layers due to efficiency and effectiveness.

| Name | Plot | Equation | Derivative |
|--|---|--|---|
| Identity |  | $f(x) = x$ | $f'(x) = 1$ |
| Binary step |  | $f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$ | $f'(x) = \begin{cases} 0 & \text{for } x \neq 0 \\ ? & \text{for } x = 0 \end{cases}$ |
| Logistic (a.k.a Soft step) |  | $f(x) = \frac{1}{1 + e^{-x}}$ | $f'(x) = f(x)(1 - f(x))$ |
| Tanh |  | $f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$ | $f'(x) = 1 - f(x)^2$ |
| ArcTan |  | $f(x) = \tan^{-1}(x)$ | $f'(x) = \frac{1}{x^2 + 1}$ |
| Rectified Linear Unit (ReLU) |  | $f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$ | $f'(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$ |
| Parametric Rectified Linear Unit (PReLU) [2] |  | $f(x) = \begin{cases} \alpha x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$ | $f'(x) = \begin{cases} \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$ |
| Exponential Linear Unit (ELU) [3] |  | $f(x) = \begin{cases} \alpha(e^x - 1) & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$ | $f'(x) = \begin{cases} f(x) + \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$ |
| SoftPlus |  | $f(x) = \log_e(1 + e^x)$ | $f'(x) = \frac{1}{1 + e^{-x}}$ |